

NLP

Tokenization

The process of dividing a text into a sequence of words.

- Different models uses different tokenization methods.

BERT TOKEN	BERT ID	GPT TOKEN	GPT ID
my	2026	My	3666
grandson	7631	Grandson	31845
loved	3866	Gloved	10140
it	2009	Git	340
!	999	G!	0
so	2061	Gso	1406
much	2172	Gmuch	881
fun	4569	Gfun	1257
!	999	G!	0

Corpus

a large and structured collection of text

- a corpus typically consists of at least a million words of text
- at least tens of thousands of distinct vocabulary words.

Text classification

- the process of categorizing text into organized groups.
- text classifiers can automatically analyze text and then assign a set of predefined tags or categories based on its content.
- machine learning approach
 - Features
 - **BoW**
 - **TF-IDF**
 - Features + classifier
 - Logistic regression
 - SVM
 - **Naive Bayes**
- Deep learning approach
 - Neural models: CNNs (capture local **n-grams**)
 - RNNs, LSTMs (sequence-aware)
 - **Transformers** (e.g. BERT)
 - **Contextual embedding**
- Modern enhancements
 - Embedding + classifier: Pretrained embeddings (**Word2Vec**) + classifier
 - Fine-tuned transformers: BERT fine-tuned

Bag of Words, BoW

- it represents a text as a bag of its words
- we can understand the meaning of a document from its content (words) their multiplicity (frequency, number of occurrences)
- mainly used as a tool of feature extraction.

- limitations
 - ignores syntax and the context
 - disregarding grammar
 - discards word order - words are independent of each other
 - considers only the meanings of the words in the sentence

```
# examples
```

```
Example1 = "He likes to watch movies. Mary likes movies too."
Example2 = "Mary also likes to watch football games."
```

```
# Vocabulary
```

```
Vocab = {"He", "likes", "to", "watch", "movies", "Mary", "also", "football", "games"}
```

```
# BoW representation
```

```
BoW1 = {"He": 1, "likes": 2, "to": 1, "watch": 1, "movies": 2, "Mary": 1, "also": 0, "football": 0, "games": 0}
```

```
BoW2 = {"He": 0, "likes": 1, "to": 1, "watch": 1, "movies": 0, "Mary": 1, "also": 1, "football": 1, "games": 1}
```

```
[[1,2,1,1,2,1,0,0,0], [0,1,1,1,0,1,1,1,1]]
```

TF-IDF

Term Frequency - Inverse Document Frequency

- Model based on the statistics of word counts.
- idea is that key terms and important ideas are likely to repeat.
- includes a scoring function that measure the relevance of a document to a query.
 - the function takes a document with a corpus and a query as input and returns a numeric score.
 - the documents that have the highest scores are considered as the most relevant documents.
- **Term Frequency**
 - $TF(q_i, d_j, D)$
- **Inverse Document Frequency**
 - $IDF(q_i, D) = \log \frac{N}{DF(q_i, D)}$
 - $DF(q_i, D)$: number of documents in the corpus D that contain the term q_i
- $N = |D|$: total number of documents in the corpus D
- **TF-IDF score**
 - $TFIDF(q_i, d_j, D) = TF(q_i, d_j) \cdot IDF(q_i, D)$

Examples of TF-IDF

- Document 1: "John likes to watch movies."
- Document 2: "Mary likes movies too."
- Document 3: "John also likes football."

Step	Term	Document	Number of times term t appears in document d	Total number of terms in document d	$TF(t, d)$	Total number of documents	Number of documents containing t
1	t=likes	d1	1	5	0.2	3	3
	t=likes	d2	1	4	0.25		
	t=likes	d3	1	4	0.25		
2	t=watch	d1	1	5	0.2	3	1
	t=watch	d2	0	4	0		
	t=watch	d3	0	4	0		
3	t=mary	d1	0	5	0	3	1
	t=mary	d2	1	4	0.25		

Step	Term	Document	Number of times term t appears in document d	Total number of terms in document d	TF(t, d)	Total number of documents	Number o containing
	t=mary	d3	0	4	0		
4	t=football	d1	0	5	0	3	1
	t=football	d2	0	4	0		
	t=football	d3	1	4	0.25		

Naive Bayes

- naive refers to a very strong simplifying assumption about the features
- Conditional independence assumption
 - Given class Y , all the features $X = X_1, X_2, \dots, X_n$ are assumed mutually independent.
- $P(X|Y) = P(X_1, \dots, X_n|Y) = \prod_{i=1}^n P(X_i|Y)$
- Bayes rule
 - $P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)} \propto P(Y)P(X|Y)$
 - $P(Y|X_1, \dots, X_n) \propto P(Y) \prod_{i=1}^n P(X_i|Y)$
- Sentiment Analysis
 - $p(C_k|x_1, \dots, x_n) = \frac{1}{Z} p(C_k) \prod_{i=1}^n p(x_i|C_k)$
 - $Z = p(x) = \sum_k p(C_k)p(x|C_k)$
 - scaling factor for normalization of $p(C_k|x)$
 - maximum a posteriori (MAP) decision rule
 - $\hat{y} = \arg \max_{k \in \{1, \dots, K\}} p(C_k) \prod_{i=1}^n p(x_i|C_k)$

Examples of Naive Bayes

$$P(\text{Class}|w_{1:N}) = \alpha \cdot P(\text{Class}) \cdot \prod_j P(w_j|\text{Class})$$

- Analyze the sentiment
 - my grandson loved it
 - $x_1 = \text{my}$, $x_2 = \text{grandson}$, $x_3 = \text{loved}$, $x_4 = \text{it}$
 - $C = \text{positive}, \text{negative}$
 - $\hat{y} = \arg \max_{k \in \{\text{positive}, \text{negative}\}} p(C_k) \cdot p(x_1|C_k) \cdot p(x_2|C_k) \cdot p(x_3|C_k) \cdot p(x_4|C_k)$
 - positive case
 - $p(\text{positive}) = 0.49$
 - $p(\text{my}|\text{positive}) = 0.30$
 - $p(\text{grandson}|\text{positive}) = 0.01$
 - $p(\text{loved}|\text{positive}) = 0.32$
 - $p(\text{it}|\text{positive}) = 0.30$
 - $p(\text{positive}|x) \propto p(\text{positive}) \cdot p(\text{my}|\text{positive}) \cdot p(\text{grandson}|\text{positive}) \cdot p(\text{loved}|\text{positive}) \cdot p(\text{it}|\text{positive}) = 0.49 \times 0.30 \times 0.01 \times 0.32 \times 0.30 = 0.00014256$
 - negative case
 - $p(\text{negative}) = 0.51$
 - $p(\text{my}|\text{negative}) = 0.20$
 - $p(\text{grandson}|\text{negative}) = 0.02$
 - $p(\text{loved}|\text{negative}) = 0.08$
 - $p(\text{it}|\text{negative}) = 0.4$
 - $p(\text{negative}|x) \propto p(\text{negative}) \cdot p(\text{my}|\text{negative}) \cdot p(\text{grandson}|\text{negative}) \cdot p(\text{loved}|\text{negative}) \cdot p(\text{it}|\text{negative}) = 0.51 \times 0.20 \times 0.02 \times 0.08 \times 0.4 = 0.00006528$
 - $Z = p(x) = \sum_k p(C_k)p(x|C_k)$
 - $= p(\text{positive}) \cdot p(\text{my}, \text{grandson}, \text{loved}, \text{it}|\text{positive}) + p(\text{negative}) \cdot p(\text{my}, \text{grandson}, \text{loved}, \text{it}|\text{negative})$
 - $= 0.00014256 + 0.00006528 = 0.00020784$
 - the probability of x , obtained by adding up its probabilities under each class.
 - $p(\text{positive}|x) = \frac{p(\text{positive}) \cdot p(\text{my}, \text{grandson}, \text{loved}, \text{it}|\text{positive})}{Z} = \frac{0.00014256}{0.00020784} \approx 0.6867$
 - $p(\text{negative}|x) = \frac{p(\text{negative}) \cdot p(\text{my}, \text{grandson}, \text{loved}, \text{it}|\text{negative})}{Z} = \frac{0.00006528}{0.00020784} \approx 0.3133$

- Decision
 - $\hat{y} = \arg \max_{k \in \{\text{positive}, \text{negative}\}} \{k = \text{positive} : 0.6867, k = \text{negative} : 0.3133\}$
 - The sentiment is **positive**.
- Spam detection
 - Probabilities learned from data:
 - $P(\text{Spam}) = 0.4, P(\text{Ham}) = 0.6$
 - $P(\text{free}|\text{Spam}) = 0.8, P(\text{free}|\text{Ham}) = 0.1$
 - $P(\text{win}|\text{Spam}) = 0.7, P(\text{win}|\text{Ham}) = 0.05$
 - New document: "free win"
 - Spam score: $0.4 \times 0.8 \times 0.7 = 0.224$
 - Ham score: $0.6 \times 0.1 \times 0.05 = 0.003$
 - Classified as **Spam**.

N-gram model

- N-gram: a sequence of written symbols of length n
 - unigram, bigram, trigram
- the probability of each symbol is dependent only on the $n-1$ previous symbols.
- $P(w_j|w_{1:j-1}) = P(w_j|w_{j-n+1:j-1})$
- $P(w_1:N) = \prod_{j=1}^N P(w_j|w_{1:j-1}) \approx \prod_{j=1}^N P(w_j|w_{j-n+1:j-1})$

Examples of N-gram

- $W_{1:N}$ is "This article is on NLP"
 - $N=5$
 - Bigram (n-gram with $n=2$)

```

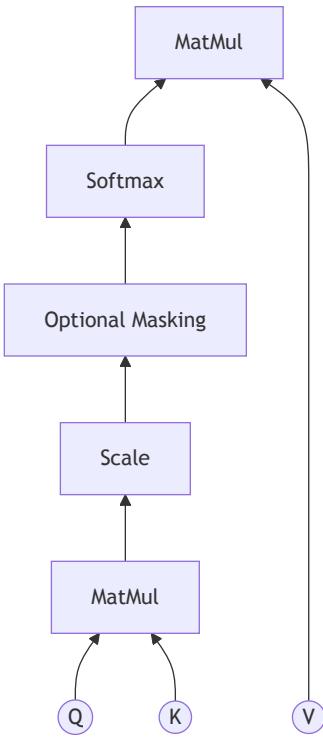
P("This article is on NLP")
/* full chain rule, */
= P("This") // j = 1
  * P("article" | "This") // j = 2
  * P("is" | "This article") // j = 3
  * P("on" | "This article is") // j = 4
  * P("NLP" | "This article is on") // j = 5
/* bigram approximation */
= P("This") // j = 1
  * P("article" | "This") // j = 2
  * P("is" | "article") // j = 3
  * P("on" | "is") // j = 4
  * P("NLP" | "on") // j = 5

```

Step (j)	Word (W_{-j})	Bigram	Trigram
1	This	$(P(\text{This}))$	$(P(\text{This}))$
2	article	$(P(\text{article} \text{This}))$	$(P(\text{article} \text{This}))$
3	is	$(P(\text{is} \text{article}))$	$(P(\text{is} \text{This, article}))$
4	on	$(P(\text{on} \text{is}))$	$(P(\text{on} \text{article, is}))$
5	NLP	$(P(\text{NLP} \text{on}))$	$(P(\text{NLP} \text{is, on}))$

Transformer

- given a set of input vectors (tokens), attention lets each token look at the others and form a weighted average of them.
- The weights are data-dependent.
- The model learns the weights representing which tokens are relevant to which other tokens.



$$X = \text{Input Embeddings}, \quad Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$

- Scores: $S = \frac{QK^T}{\sqrt{d}}$
 - dot products between every query and every key
 - the scale \sqrt{d} keeps gradients stable
- Weights: $A = \text{Softmax}(S)$
 - row-wise softmax
 - each row sums to 1
- Output: $\text{Attention}(Q, K, V) = AV$
 - each output token is a weighted sum of the value vectors, with weights determined by the attention scores.(how much to pay attention to each position)
- Query (Q): What am I looking for?
- Key (K): What is the label/address of the information I have?
- Value (V): What is the actual information I want to convey?

Cross-Attention

- look up relevant information in another sequence.
 - e.g. decoder attending to encoder outputs in translation
- Q from the current sequence, K and V from the other sequence.
 - $Q = X_{\text{target}}W_Q, \quad K = X_{\text{source}}W_K, \quad V = X_{\text{source}}W_V$
- Cross-attention = Which source tokens are relevant to this target token?
- Self-attention = Which other tokens in this sequence are relevant to this token?
- $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$
- $\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W_O$

Word representation

One-hot representation

- represents each word as a vector of length N .
- where N is the size of the vocabulary.

```
vocab = {he, is, singing, she, dancing, stage}
```

```
she = [0, 0, 0, 1, 0, 0]
is = [0, 1, 0, 0, 0, 0]
singing = [0, 0, 1, 0, 0, 0]
```

- limitations
 - incredibly inefficient for large vocabularies
 - not embed any intrinsic meaning of words
 - unable to represent similarity between likely words
 - the representation of documents is sparse vectors
 - can cause challenges in computation

Word Embedding

- represents individual words as vectors in a low-dimensional continuous space.
- a distributed representation of a word.
- generate a unique value for each word while using smaller vectors compared with one-hot encoding.
- common vector dictionaries
 - word2vec
 - Glove (Global Vectors)

Contextual embedding

- generates different vectors for the same word based on its context.
- the word "bank" would have different embeddings in the sentences:
 - "He went to the bank to deposit money."
 - "She sat by the river bank and enjoyed the view."
- BERT or GPT use deep neural networks to process a sequence of tokens.
- Each token's embedding is computed by considering the token itself, its position in the sequence and the surrounding tokens, context.
- captures both semantic and syntactic role in that specific context.

Part of Speech (POS) tagging

- lexical category or tag that indicates the grammatical role of a word in a sentence.
- parts of speech allow language models to capture generalizations such as "adjectives often modify nouns" or "verbs often follow subjects".

```
From the start , it took a person
IN DT NN , PRP VBD DT NN
```

```
with great qualities to succeed
IN JJ NNS T0 VB
```

Tag	Description	Example
CC	Coordinating conjunction	and, but
CD	Cardinal number	one, two
DT	Determiner	the, a
EX	Existential there	there
FW	Foreign word	doppelgänger
IN	Preposition or subordinating conjunction	in, of
JJ	Adjective	big, old
JJR	Adjective, comparative	bigger, older
JJS	Adjective, superlative	biggest, oldest

Tag	Description	Example
LS	List item marker	1, 2, One
MD	Modal	can, will
NN	Noun, singular or mass	cat, tree
NNS	Noun, plural	cats, trees
NNP	Proper noun, singular	John, London
NNPS	Proper noun, plural	Smiths, Londons
PDT	Predeterminer	all, both
POS	Possessive ending	's, s'
PRP	Personal pronoun	I, you, he
PRP\$	Possessive pronoun	my, your, his
RB	Adverb	quickly, very
RBR	Adverb, comparative	faster, better
RBS	Adverb, superlative	fastest, best
RP	Particle	up, off
SYM	Symbol	\$, %, &
TO	to	to
UH	Interjection	oh, wow
VB	Verb, base form	be, have
VBD	Verb, past tense	was, had
VBG	Verb, gerund or present participle	being, having
VBN	Verb, past participle	been, had
VBP	Verb, non-3rd person singular present	talk, have
VBZ	Verb, 3rd person singular present	talks, has
WDT	Wh-determiner	which, that
WP	Wh-pronoun	who, what
WP\$	Possessive wh-pronoun	whose
WRB	Wh-adverb	where, when
#	Pound sign	#
\$	Dollar sign	\$
,	Comma	,
.	Sentence-final punctuation	. ! ?

Example of POS tagging

- Hidden Markov Model (HMM)
 - takes in a temporal sequence of evidence observations
 - predicts the lexical categories
- Logistic regression
 - build 45 different logistic regression models, one for each part of speech

- ask each model how probable it is that the example word is a member of that category, given the feature values for that word in its particular context.

Machine translation

- translate a sentence from a source language to a target language.
- train an MT model: a large corpus of source/target sentence pairs and hope that the trained MT model can accurately translate new sentences.
- want to generate a target language sentence that corresponds to the source language sentence
- the generation of each target word is conditional on the entire source sentence and on all previously generated target words.

Example of machine translation

- a sequence-to-sequence model
 - use two RNNs (LSTM)
- attentional sequence-to-sequence model
 - use attention to create a context-based summarization of the source sentence into a fixed-dimension representation
- transformer-based model
 - encoder: reads the source sentence and turns it into a rich, contextual set of vectors.
 - decoder: generates the target sentence one token at a time, using what it has generated so far and the encoder's representations.

Text generation

- a subfield of NLP
- leverages knowledge in computational linguistics and AI to automatically generate natural language texts
- can satisfy certain communicative requirements

Example of text generation

- Classifier based on word embeddings, e.g. RNN and LSTM
 - RNN: each input word is encoded as a word embedding vector x_i , a hidden layer z_t , the classes are the words of the vocabulary
 - the output y_t will be a softmax probability distribution over the possible values of the next word in the sentence.
 - LSTM: can choose to remember some parts of the input, copying it over to the next time step, and to forget other parts.
- Pre-trained language model using deep learning
 - BERT
 - GPT-X, Generative Pre-trained Transformer

Transfer learning

- experience with one learning task helps an agent learn better on another task.
- pretraining: a form of transfer learning in which we use a large amount of shared general-domain language data to train an initial version of an NLP model.
 - we can use a smaller amount of domain-specific data to refine the model
 - the refined model can learn the vocabulary, idioms, syntactic structure, and other linguistic phenomena that are specific to the new domain.
- For NN, learning consists of adjusting weights, so the most plausible approach for transfer learning is to copy over the weights learned for task A to a network that will be trained for task B.
 - The weights are then updated by gradient descent in the usual way using data for task B.
- the popularity of transfer learning is the availability of high-quality pretrained models.
- will want to freeze the first few layers of the pretrained model
 - these layers serve as feature detectors that will be useful for new model.
 - new data set will be allowed to modify the parameters of the higher levels only
 - these are the layers that identify problem-specific features and do classification.