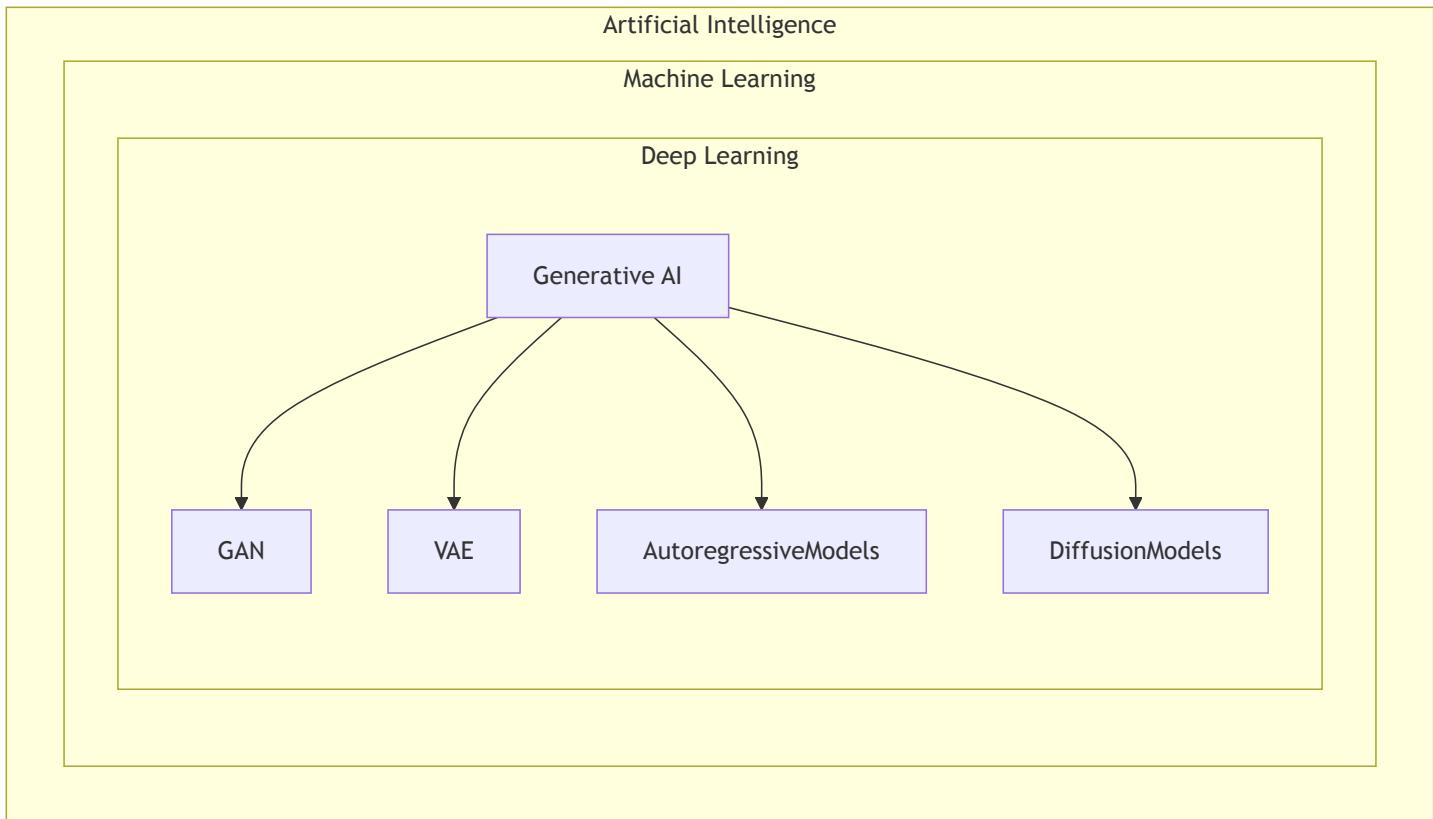


Generative AI

- Gen AI refers to a category of AI models designed to generate new content, synthetic data that resembles a given dataset.
- Gen AI models create new content, including text, images, audio, and video.



Historical context of Gen AI

- 1980s: The development of statistical approaches to AI emerged, focusing on probabilistic models
- 1990s: Hidden Markov Models (HMMs) became popular in speech recognition and sequence generation tasks, making a shift towards using statistical methods in generative processes.
- 1990s-2000s: The resurgence of Neural Networks, with the introduction of deep learning techniques. However, hardware and data limitations hampered progress.
- 2010s: The advent of deep learning algorithms, especially convolutional neural networks (CNNs) and recurrent neural networks (RNNs), transformed the landscape of AI.
- 2013: VAEs were introduced, providing a probabilistic approach to data generation, allowing for smooth latent space interpolation and structure.
- 2014: GANs, a novel framework where two neural networks (a generator and a discriminator) are trained simultaneously, allowing for the generation of highly realistic images and other data

types.

- 2015-2020: Generative models began to find applications beyond image synthesis, including text generation, music composition, and even video generation. OpenAI's GPT-2 showcased the potential of transformer-based models for generating coherent text.
- 2020s: The introduction of larger and more capable models like OpenAI's GPT-3 and subsequent iterations revolutionized natural language processing.
 - DALL-E and Stable Diffusion pushed the boundaries of image generation, allowing users to create images from textual descriptions. This sparked creative exploration and practical applications in marketing, art, and design.
- Present: The integration of multiple data modalities (text, image, audio) led to the development of models like CLIP and GPT-4, which can understand and generate content across different formats, enhancing the versatility of generative AI.

Variational Autoencoders, VAEs

- VAEs are from the probabilistic approach to build a generative AI model
- VAE learns a latent distribution instead of a fixed latent presentation, allowing for the generation of new samples.
- Latent space
 - sample from a Gaussian distribution
 - μ and σ are learned during training
 - z is a new sample from the latent space
- $q_\phi(z|x)$ is the encoder that maps input data x to a distribution over the latent space z
- $p_\theta(x|z)$ is the decoder that maps a latent variable z back to the data space x
- During training, the model optimizes the reconstruction loss and a regularization term (KL divergence) to ensure the learned latent distribution is close to a prior distribution (usually a standard normal distribution).
- Latent Representation
 - Dimensionality Reduction: latent space is typically lower-dimensional than the input space
 - Smoothness: In a well-structured latent space, similar inputs will be represented by nearby points.
 - Generative Capabilities: Once trained, can sample from the latent space to generate new data that resembles the training set.
 - Regularization: The VAE incorporates a regularization term in its loss function (KL divergence), which encourages the learned latent distribution to be close to a standard normal distribution.
- Decoder

- The decoder takes the sampled latent representation z and reconstructs the original input data.
- The goal is to make the reconstructed data as close as possible to the original input.

VAE Examples

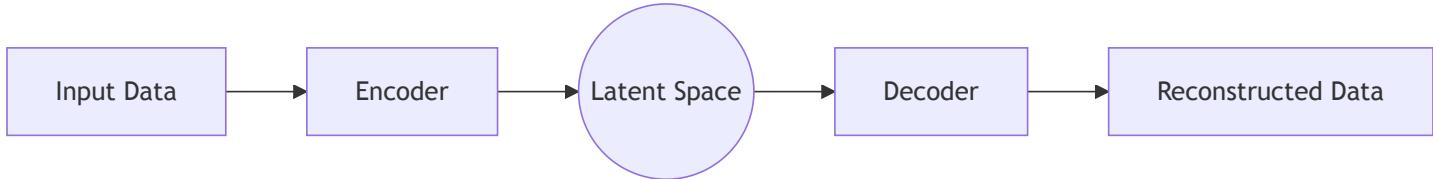
Face Generation Example

- Trains a VAE on face photos.
- Latent space learns meaningful features (e.g. hair color, glasses, smile)
- By sampling and interpolating in latent space, can generate new faces or smoothly morph one face into another.

Anomaly Detection Example

- Train a VAE on normal sensor data.
- When fed unusual data, reconstruction error will be high.
- Use this for fault detection in machines, fraud detection, etc.

Autoencoder



- Encoder
 - Input Layer: takes in the original data
 - Hidden Layer: These layers progressively reduce the dimensionality of the input through operations like linear transformations and non-linear activations. (e.g., ReLU-Rectified Linear Unit)
 - Output Layer: Produces the final latent representation. Sometimes Sigmoid or Tanh are used depending on the nature of the input data.
 - Learns to capture meaningful patterns in the data by minimizing reconstruction loss, which measures the difference between the original input and the reconstructed output produced by the decoder.
- Decoder
 - typically mirrors the structure of the encoder but in reverse.
 - During training, Mean Squared Error (MSE) for continuous data or Binary Cross-Entropy (BCE) for binary data are commonly used to quantify the reconstruction loss.

Image Reconstruction Example

- input: 28x28 pixel
- encoder: compresses it to just 16 numbers (latent vector)
- decoder: expands those 16 numbers back into a 28x28 image
- result: the reconstructed digit looks similar to the original but not in the training set.

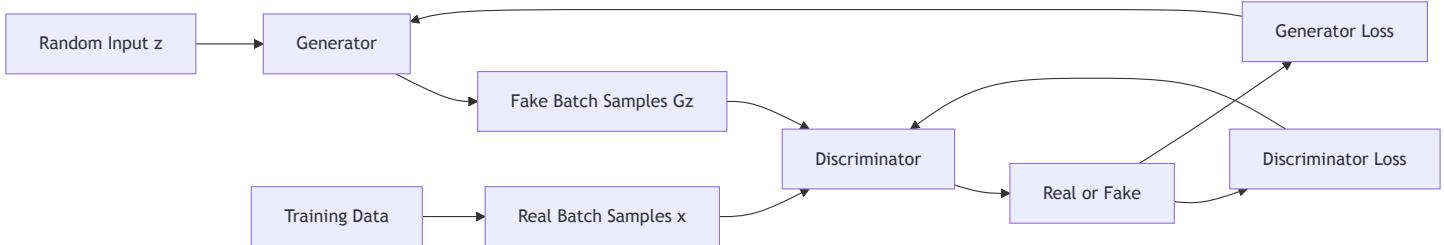
Denoising Images

- input: a noisy photo of a cat
- encoder: learns to ignore the noise and compress meaningful features.
- decoder: rebuilds the image without the noise.
- output: a clearer cat image.

Loss function

- Reconstruction loss: Ensure output similar to input
- KL divergence loss: Push the learned distribution to be close to a standard normal distribution.
(can sample new data)

Generative Adversarial Network, GAN



- Generator: Learns to generate fake (generated) data that resembles data distribution.
- Discriminator: Learns to distinguish between real data and data generated by the Generator.
- Steps
 - i. Generator creates fake data samples from random input (noise).
 - ii. Discriminator evaluates these samples together with real data samples.
 - iii. Discriminator outputs probabilities indicating whether each sample is real or fake.
 - iv. Based on the prediction of the Discriminator, the Generator and Discriminator are updated using specific loss functions.
- Applications
 - Labels to Street Scenes
 - Labels to Facade

- BW to Color
- Aerial to Map
- Day to Nihght
- Edges to Photo
- Text-to-Image Synthesis

Autoregressive Models

- a class of generative models where the current value of a time series is expressed as a linear function of its own past values plus some noise.
- foundational in generative AI and widely used in generative AI, particularly for tasks like text generation, speech synthesis.
- Examples
 - GPT series: state-of-the-art autoregressive language models used for text generation.
 - WaveNet: an autoregressive model for generating high-quality audio by predicting each sample conditioned on previous samples.

Training Autoregressive Models

- Pretraining
 - the model sees sequences of tokens and learns to guess the next token.
 - objective: minimize cross-entropy loss between its guess and the true next token.
 - this teaches grammar, facts, reasoning patterns, and style from raw text.
- Supervised fine-tuning (optional)
 - smaller curated datasets (questions to answer, instructions to respond) teach it to follow directions.
- Reinforcement learning from Human Feedback (optional)
 - Humans rank multiple model outputs; a reward model is trained on those rankings; GPT is then optimized to produce higher-reward responses (safer, more helpful, less toxic)
- Text to **Tokens** via a tokenizer (e.g. BPE, Byte Pair Encoding)
- Each token becomes a vector (**Embedding**)
- **Positional encodings** inject order information.
- **Transformer layer**
 - **self-attention**: each token looks at all previous token (causal mask) and decides which ones matter, computing weighted combinations.
 - Multiple **heads** let it attend to different patterns (syntax, long-range links, etc.).
 - **feed-forward network**: a nonlinear MLP refines each token's representation.
 - **residual connections & layer norm** stabilize training.

Interference

- a prompt (the context)
- GPT computes probabilities for the next token.
- A decoding strategy samples a token
 - Decoding knobs
 - Greedy: take the top token
 - Top-K
 - Nucleus (Top-p) sampling (limit to likely options)
 - Temperature scales randomness, controls how random the next token choice is
 - lower = more deterministic, higher = more creative
- append the token and repeat autoregressive generation.